

Bounding tree-based Sparse PCA using perturbation theory

Daniel Ringwalt

January 5, 2023

1 Perturbing the partial solution

- Some nodes in the tree will commit an additional variable to the partial solution.
- The node can be represented by an indicator vector $z \in \{0, 1\}^n$ having lower and upper bounds.
- The nodes of interest will assign a value of 1 at a formerly unbounded location.
- Sparse PCA is given by the SVD of a completed data subset M_z , $z \in \{0, 1\}^n, \|z\|_1 = k$.

2 Lower bound from the partial solution

- Use the SVD of the partially built data matrix M_1 .
- Specifically, we can compute the dominant left singular vector u .
- There is a different u after subsetting M according to the global optimum k -sparse PCA solution.
- Suppose that the correlation of u and the globally optimum u is very high.
- Select the remaining $k - i$ variables from M greedily according to their dot product with u .
- We may find a better u and objective value on this greedy subset by running SVD again, or assume that u is of high quality and take the sum of dot products as a lower bound on the objective.
- This deterministic (over some fixed M_1) algorithm will be evaluated against stochastic sparse PCA.

3 Upper bound on completing the solution

- Completing M_1 is a combinatoric problem to be addressed in the next section.
- Suppose that M_1 is a good approximation of the i -sparse PCA problem, so the left-SVD of M_1 is starting to approach the left-SVD of our solution M_z .
- Consider the Gram matrix $M_1 M_1^T$ (where the entries of the matrix represent the observations): Σ_1 .
- The matrix Σ_1 has an eigensystem where only the dominant eigenvalue is of interest, so we can approximate Σ_1 as being rank-one.
- We can complete Σ_1 by adding a rank- $k - i$ matrix Σ_{k-i} . Assume that this matrix is fixed, and consider the combinatorial problem in the next section.
- The operator norm, or definite matrix spectral radius, $\|\Sigma_1 + \Sigma_{k-i}\|$ is upper-bounded by $\|\Sigma_1\| + \|\Sigma_{k-i}\|$.
- Project the data onto u : $\|\Sigma_{k-i}\|_2 = \|M_{k-i}^T u\|_2^2 + \|\Sigma_{k-i}^{\text{res}}\|_2$.
- The operator norm $\|\Sigma_{k-i}^{\text{res}}\|_2$ is given by some vector $v, \|v\|_2 = 1$ which maximizes $\|\Sigma_{k-i}^{\text{res}} v\|_2$.
- In this case, we constructed $\Sigma_{k-i}^{\text{res}}$ to be nearly singular by subtracting the projection onto u , so u and this dominant v are almost orthogonal. The bound is not very tight.
- For perturbing one positive semidefinite eigensystem by adding another such matrix, we can start from the dominant eigenvalue/eigenvector of one particular summand, and use a matrix-vector multiplication for the other summand:

$$\|\Sigma_1 + \Sigma_{k-i}\| \leq \|M_1 u\|_2^2 + \|M_{k-i} u\|_2^2 + \|\Sigma_{k-i}^{\text{res}} u\|_2$$

4 Parameterizing a solution

Revisit the indicator vector $z \in \{0, 1\}^{n-i}$, $\|z\|_1 = k - i$ (the contribution of the original i variables to the objective function has already been made constant). The objective function should be heavily influenced by the rank-one projection, with a greater contribution than the residual's contribution: $\sum_i z_i (M_{*i}^T u)^2$.

4.1 Residual term

For the residual's contribution, we need to insert a sparsity-inducing term into the middle of the matrix-vector multiplication over the full covariance matrix: $\|\Sigma_{n-i}^{\text{res}} v\|_2$. We sparsify the rows of Σ and of u using a diagonal matrix on the left side, but the sparsified u should be unit-normed:

$$\frac{\|\text{Diag}(z)\Sigma_{n-i}^{\text{res}}\text{Diag}(z)v\|}{\|\text{Diag}(z)v\|}$$

4.2 Limitations of quadratic programming

Our objective function could start to look like a matrix-vector multiplication quadratic. Just like PCA (which is solved by the dominant eigenvector), the equation $\max v^T \Sigma v$ cannot be solved via an optimization program. Usually, convex optimization solvers would be applied to a positive definite matrix of costs/risks to be minimized, not maximized.

Plug-ins can be created for some optimization libraries in the form of custom cones. For example, the set $(t \in \mathbb{R}, x \in \mathbb{R}^n) : t \geq \|x\|_2$ is a quadratic cone constraint. The cone must be closed under addition (in $(t_1, x_1) + (t_2, x_2)$, the t grows very quickly and the bound clearly holds in this direction). Also, the cone must be closed under scaling by a real number (so we cannot bound the reciprocal ℓ_2 norm in the direction which we might want, for scaling our objective function post-hoc).

4.3 Semidefinite programming relaxation

Rewriting our objective function using a rank-one zero-one matrix zz^T helps us find $\|z\|_2^2$ using $\text{Tr}(zz^T)$. A rank-one constraint would not be a conic constraint (the set of matrices is clearly not closed under addition), so the constraint is usually dropped. Dropping $\text{Diag}(z)$, use a PSD matrix (expected to be low-rank) S (sparsity-inducing matrix).

$$\|\text{Diag}(z)v\|_2^2 = v^T \text{Diag}^2(z)v \approx v^T S v = \text{Tr } v^T S v = \text{Tr } S v v^T = \langle S, v v^T \rangle$$

The sparsity-inducing matrix S will be applied to our residual matrix here, so use a scaled matrix S_{res} with constraint (v defined ahead of time) $\langle S_{\text{res}}, v v^T \rangle = 1$. Later, we will scale S back up by some variable scale factor (which cannot be expressed ahead of time in a linear program).

For the (squared) numerator, we will need to generalize this inner product to encompass 3 PSD matrices. The inner product $\langle A, B \rangle$ can be expressed as $\sum_i \sum_j a_{ij} b_{ij} = \sum_i \sum_j (A \odot B)_{ij}$, and we will generalize this to applying 3 PSD matrices.

$$\|\text{Diag}(z)\Sigma_{n-i}^{\text{res}}\text{Diag}(z)v\|_2^2 \approx \|(S_{\text{res}} \odot \Sigma_{n-i}^{\text{res}})v\|_2^2 = v^T (S_{\text{res}} \odot \Sigma_{n-i}^{\text{res}})^2 v$$

In general, introducing square convex terms into a maximization (not minimization) problem would lead to NP-hardness (a la quadratic programming minimization with an indefinite matrix). We cannot square the matrix, and will bound the action of the matrix's spectrum on this vector v . First, we have $vS_{\text{res}}v^T = 1$ as one of our constraints. Elementwise multiplication by the matrix $\Sigma_{n-i}^{\text{res}}$ produces another positive semidefinite matrix (Schur product theorem). This matrix can be loosely bounded using the largest absolute value element of Σ , which is found on its diagonal, multiplied by S_{res} . For any bound B on $vMv^T \leq B$ for positive semidefinite M , we expect $vM^2v^T \leq vMv^T * B$.

$$\|\text{Diag}(z)\Sigma_{n-i}^{\text{res}}\text{Diag}(z)v\|_2^2 < \langle S_{\text{res}}, \Sigma_{n-i}^{\text{res}} \odot vv^T \rangle * \max_i |(\Sigma_{n-i}^{\text{res}})_{ii}|$$

Finally, this term in the objective requires a new scalar value which is the square root. The square root function is concave. The quadratic cone can allow for slack in the inequality $|a| \leq \sqrt{b}$, so conic programming is useful for a square root term which represents a valuable quantity, but not for representing the square root of a cost.

4.4 Rank-one projection term

The impact of approximations above is limited in scope, since the problem should have reasonable eigenvalue separation and the current rank-one principal component should be a guess at diagonalizing the problem. Therefore, we add large contributions to the problem (rank-one projections of each variable) which are linear in the diagonal entries of S . If the positive coefficients in the entries of S dominate the problem, then the problem reduces to selecting $S = zz^T, z \in \{0, 1\}^n, \|z\|_0 = k$.

4.5 Scaling S using the power cone

If S_{res} is rank-one, then we can recover the vector representing its range, applying some scale factor to be determined. We simply need a square root and arbitrary scaling of the diagonal entries of S_{res} (use positive square root, since all entries of S_{res} are constrained to be nonnegative). This can be done in conic programming using the power cone $\mathcal{P}_3^{1/2}$, which is the set of triples $(x_1, x_2, x_3) : \sqrt{x_1x_2} \geq |x_3|$. We will have n conic constraints where x_2 is some scalar variable not otherwise used, x_1 is a diagonal entry of S_{res} , and x_3 is a problem variable taking the rank-one projection as a coefficient.

The free variable x_2 in the conic constraints will generally be greater than 1, to scale up the sparsity S entries to a unit value. The x_3 problem variables (z_i) are constrained by the unit box. The z_i variables are used to sum up the contribution of variance from the rank-one projection, and should be at unity.